

Short Note

The most important decision I made was keeping the LLM swappable instead of hardwiring one provider into the code.

`classify()` and `synthesize()` sit behind a fixed contract, the same inputs, outputs, and prompts no matter what's underneath, and all the provider-specific code lives in one file. That turned out to matter more than I expected. I started on Anthropic, hit a billing wall, tried OpenAI and hit the same problem, and finally moved to Google Gemini's free tier by editing a single file. Nothing else changed. The rules, the pipeline, and every function signature stayed exactly the same.

The trade-off was the free tier's limits. It only allows about 20 requests a day, which at first barely covered a single insight. Instead of paying to get around it, I batched the comment classification into one call, which took a full insight from around 41 requests down to 2. That made the free tier workable and the whole thing cheaper to run.

So the cost is leaning on a rate-limited free tier and an SDK that's now legacy, both noted in the README with how I'd fix them in production. I'd rather the system survive a provider going down than be tied to one.